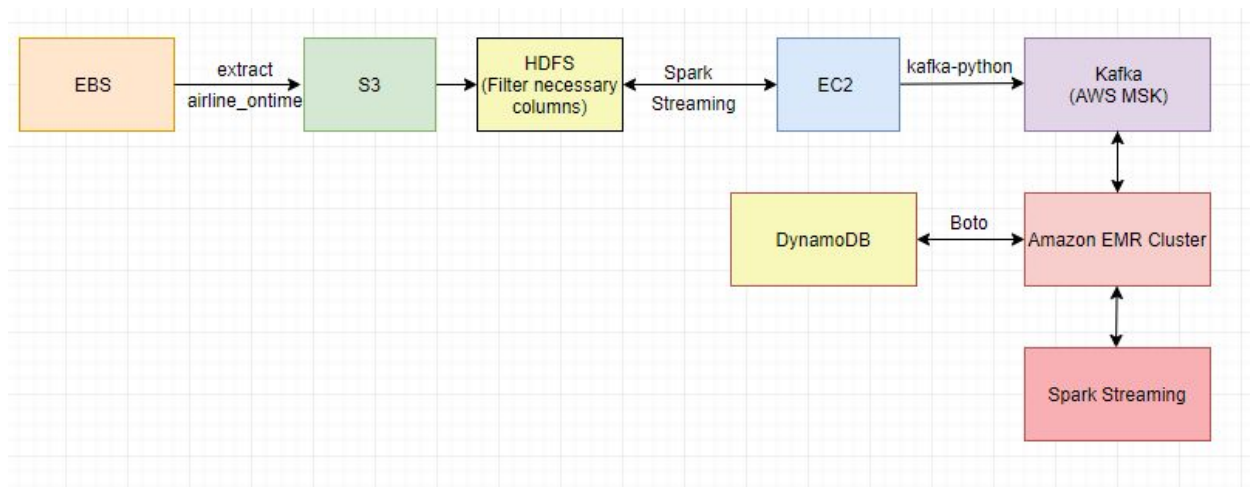


Overall Design



- For Task 2, I have adopted Kafka + Spark Streaming
- Airline_ontime folder is unzipped to S3.
- Relevant columns from CSV are filtered and stored in HDFS which is then streamed to EC2. EC2 instance then writes the data to Kafka topic
- AWS MSK is used to set up a Kafka Cluster with replication factor = 3 one m5.x2large node in each of the 3 regions. Kafka version used in 2.1.0
- Data in Kafka topic is consumed by EMR cluster, task2 queries are processed and the results stored in DynamoDB using boto python library.

Data Extraction

EBS Volume is mounted on EC2 and then only the required folder (airline_ontime) is unzipped to S3.

```
./ebs_to_s3.sh $EBS_MOUNTED_FOLDER $S3_BUCKET
```

Source

[Extraction Script](#)

Data in S3 is copied to HDFS and relevant columns filtered using spark streaming

Source

[streamFiltering.py](#) (discards cancelled flights)

[streamFilteringWithCancelled.py](#) (includes cancelled flights)

[streamFiltering_2008.py](#) (includes info only for 2008 flights and discards cancelled flights)

- Filtered data is moved to kafka cluster under different topics. **Kafka-python** library is used for this

Some optimization flags used to prevent data drop and speed up write

```
batch_size=98304 (default batch size is 16k)
linger_ms=100
acks='all' (acks set to false by default)
```

Source

[kafkaProducer.py](#)

[kafkaProducer_2008.py](#)

Snapshot of kafka topics created

```
root@ip-172-31-0-210 kafka_2.12-2.1.0# bin/kafka-run-class.sh kafka.tools.GetOffsetShell --broker-list b-2.kafkacluster.qa2zr3.c2.kafka.us-east-1.amazonaws.com:9092,b-3.kafkacluster.qa2zr3.c2.kafka.us-east-1.amazonaws.com:9092,b-1.kafkacluster.qa2zr3.c2.kafka.us-east-1.amazonaws.com:9092 --topic airportsAll12 --time -1 --offsets 1 | awk -F ":" '{sum += $3} END {print sum}'
5824436
root@ip-172-31-0-210 kafka_2.12-2.1.0# bin/kafka-topics.sh --zookeeper 172.31.0.135:2181,172.31.1.124:2181,172.31.2.25:2181 --list
airportsAll12
airportsFull
airportsWithCancelled
airportsWithCancelled2
```

Number of messages in kafka topic for year 2008 (not including cancelled flights)

5824436

```
root@ip-172-31-0-210 kafka_2.12-2.1.0# bin/kafka-run-class.sh kafka.tools.GetOffsetShell --broker-list b-2.kafkacluster.qa2zr3.c2.kafka.us-east-1.amazonaws.com:9092,b-3.kafkacluster.qa2zr3.c2.kafka.us-east-1.amazonaws.com:9092,b-1.kafkacluster.qa2zr3.c2.kafka.us-east-1.amazonaws.com:9092 --topic airportsAll12 --time -1 --offsets 1 | awk -F ":" '{sum += $3} END {print sum}'
5824436
```

Number of messages in kafka topic for all years (not including cancelled flights)

114345531

```
root@ip-172-31-0-210 kafka_2.12-2.1.0# bin/kafka-run-class.sh kafka.tools.GetOffsetShell --broker-list b-2.kafkacluster.qa2zr3.c2.kafka.us-east-1.amazonaws.com:9092,b-3.kafkacluster.qa2zr3.c2.kafka.us-east-1.amazonaws.com:9092,b-1.kafkacluster.qa2zr3.c2.kafka.us-east-1.amazonaws.com:9092 --topic airportsFull --time -1 --offsets 1 | awk -F ":" '{sum += $3} END {print sum}'
114345531
```

Number of messages in kafka topic for all years (including cancelled flights)

116754192

```
root@ip-172-31-0-210 kafka_2.12-2.1.0# bin/kafka-run-class.sh kafka.tools.GetOffsetShell --broker-list b-2.kafkacluster.qa2zr3.c2.kafka.us-east-1.amazonaws.com:9092,b-3.kafkacluster.qa2zr3.c2.kafka.us-east-1.amazonaws.com:9092,b-1.kafkacluster.qa2zr3.c2.kafka.us-east-1.amazonaws.com:9092 --topic airportsWithCancelled2 --time -1 --offsets 1 | awk -F ":" '{sum += $3} END {print sum}'
116754192
```

- Kafka Consumer is EMR cluster with Spark 2.4.2 which reads from the topics and stores final answers in dynaodb

Problem 1.1

Strategy:

1. map() each airport_id to (airport_id,1)
2. reduce() count for each airport_id to get (airport_id,count)
3. filter out top 10 airports with highest counts

```
spark-submit --executor-memory 6g --packages  
org.apache.spark:spark-streaming-kafka-0-8_2.11:2.1.0 --conf  
spark.streaming.backpressure.enabled=true --conf  
spark.streaming.kafka.maxRatePerPartition=250000 g1_1.py
```

Source:

[g1_1.py](#)

Results:

Format: (AirportID, Count)

```
19/06/04 06:46:56 INFO DStreamGraph: Updated checkpoint data for time 1559628726000 ms  
19/06/04 06:46:56 INFO CheckpointWriter: Submitted checkpoint of time 1559628726000 ms to writer queue  
19/06/04 06:46:56 INFO CheckpointWriter: Saving checkpoint for time 1559628726000 ms to file 's3://mudabircapstonecheckpoint/top10airports/checkpoint-1559628726000'  
(u'ORD', 12446097)  
(u'ATL', 11537401)  
(u'DFW', 10795494)  
(u'LAX', 7721141)  
(u'PHX', 6582467)  
(u'DEN', 6270420)  
(u'DTW', 5635421)  
(u'IAH', 5478257)  
(u'MSP', 5197649)  
(u'SFO', 5168898)  
19/06/04 06:46:57 INFO JobScheduler: Finished job streaming job 1559629098000 ms.0 from job set of time 1559629098000 ms
```

Problem 1.2

Strategy

- 1) Filter for all non cancelled flights and map (flightID, ArrivalDelay)
- 2) Get (flightID, (delaySum, count, avgArrivalDelay) //see updateStateByKey
- 3) Filter top 10 flights with best Arrival delay

Source:

[g1_2.py](#)

Results:

Format: FlightID [totalDepDelay, totalCount, AvgArrivalDelay]

```
(u'HA', (-264258.0, 261175, -1.01180434574519))
(u'AQ', (175282.0, 151507, 1.1569234424812056))
(u'PS', (60319.0, 41581, 1.4506385127822803))
(u'ML (1)', (328150.0, 69119, 4.747609195734892))
(u'PA (1)', (1570649.0, 293971, 5.342870555258852))
(u'F9', (1751640.0, 320468, 5.465881148819851))
(u'NW', (52626818.0, 9468157, 5.558295875321882))
(u'WN', (82237109.0, 14794127, 5.558767272986098))
(u'OO', (16842083.0, 2936047, 5.736312463662878))
(u'9E', (2713127.0, 462424, 5.8671846616957595))
```

```
19/06/08 09:37:05 INFO JobScheduler: Added jobs for time 1559985912000 ms
19/06/08 09:37:05 INFO JobScheduler: Starting job streaming job 1559985912000 ms.0 from job set of time 1559985912000 ms
(u'HA', (-264258.0, 261175, -1.01180434574519))
(u'AQ', (175282.0, 151507, 1.1569234424812056))
(u'PS', (60319.0, 41581, 1.4506385127822803))
(u'ML (1)', (328150.0, 69119, 4.747609195734892))
(u'PA (1)', (1570649.0, 293971, 5.342870555258852))
(u'F9', (1751640.0, 320468, 5.465881148819851))
(u'NW', (52626818.0, 9468157, 5.558295875321882))
(u'WN', (82237109.0, 14794127, 5.558767272986098))
(u'OO', (16842083.0, 2936047, 5.736312463662878))
(u'9E', (2713127.0, 462424, 5.8671846616957595))
19/06/08 09:37:05 INFO JobScheduler: Finished job streaming job 1559985912000 ms.0 from job set of time 1559985912000 ms
19/06/08 09:37:05 INFO JobScheduler: Total delay: 713.794 s for time 1559985912000 ms (execution: 0.132 s)
```

Problem 2.1

```
'''
```

The incoming data format is

```
Year|Month|date|DayofWeek|UniqueCarrier|FlightNum|Origin|Dest|CRSDeptime|DepD
elay|ArrDelay
'''
```

(For each origin airport X, rank top 10 carriers in decreasing order of ontime departure)

Strategy:

- 1) Filter for all non cancelled flights and map ((OriginAirport,carrier), DepDelay)
- 2) Get ((OriginAirpor,carrier), (depDelaySum,count, avgDepDelay)) // see
`flightsDelay.updateStateByKey(updateFunction)`
- 3) For each OriginAirport, filter 10 best carriers in terms of Dep delay
- 4) Store results in dynamoDB
 - Main DB has OriginAirport as Partition key and Carrier as Sort key
 - Also created a Secondary index mapped to main DB with Partition key as OriginAirport and sort key as DepDelay. Hence querying the Secondary index returns results sorted in increasing order of DepDelay.

Source:

[g2_1.py](#)

Results:

Query: [Table] Top10CarriersTask2: Origin, Carrier ▾			
<input type="checkbox"/>	Origin ▾	Carrier ▾	DepDelay ⓘ ▲
<input type="checkbox"/>	SRQ	TZ	-0.381996974281
<input type="checkbox"/>	SRQ	XE	1.48976677772
<input type="checkbox"/>	SRQ	YV	3.40402193784
<input type="checkbox"/>	SRQ	AA	3.58326653307
<input type="checkbox"/>	SRQ	UA	3.95212206243
<input type="checkbox"/>	SRQ	US	3.96839828967
<input type="checkbox"/>	SRQ	TW	4.30467606502
<input type="checkbox"/>	SRQ	NW	4.85635924135
<input type="checkbox"/>	SRQ	DL	4.86917943416
<input type="checkbox"/>	SRQ	MQ	5.35058823529

Query: [Table] Top10CarriersTask2: Origin, Carrier ▾			
<input type="checkbox"/>	Origin ▾	Carrier ▾	DepDelay ⓘ ▲
<input type="checkbox"/>	CMH	DH	3.49111470113
<input type="checkbox"/>	CMH	AA	3.51526494895
<input type="checkbox"/>	CMH	NW	4.04155500526
<input type="checkbox"/>	CMH	ML (1)	4.36645962733
<input type="checkbox"/>	CMH	DL	4.71344133974
<input type="checkbox"/>	CMH	PI	5.20129487934
<input type="checkbox"/>	CMH	EA	5.93738938053
<input type="checkbox"/>	CMH	US	5.99220342136
<input type="checkbox"/>	CMH	TW	6.15909742531
<input type="checkbox"/>	CMH	YV	7.96119133574

Query: [Index] CarriersByDepDelayTask2: Origin, DepDelay ▾				
<input type="checkbox"/>	Origin ▾	Carrier ▾	DepDelay ▾	
<input type="checkbox"/>	JFK	UA	5.96832536487	
<input type="checkbox"/>	JFK	XE	8.11373626374	
<input type="checkbox"/>	JFK	CO	8.20120808165	
<input type="checkbox"/>	JFK	DH	8.74298090807	
<input type="checkbox"/>	JFK	AA	10.0978800018	
<input type="checkbox"/>	JFK	B6	11.1270962227	
<input type="checkbox"/>	JFK	PA (1)	11.5555935761	
<input type="checkbox"/>	JFK	NW	11.6378177165	
<input type="checkbox"/>	JFK	DL	11.98453575	
<input type="checkbox"/>	JFK	TW	12.641537803	

Query: [Index] CarriersByDepDelayTask2: Origin, DepDelay ▾				
<input type="checkbox"/>	Origin ▾	Carrier ▾	DepDelay ▾	
<input type="checkbox"/>	SEA	OO	2.70581965466	
<input type="checkbox"/>	SEA	PS	4.72063933287	
<input type="checkbox"/>	SEA	YV	5.12226277372	
<input type="checkbox"/>	SEA	TZ	6.34500393391	
<input type="checkbox"/>	SEA	US	6.41238418226	
<input type="checkbox"/>	SEA	NW	6.49876240739	
<input type="checkbox"/>	SEA	DL	6.5280229259	
<input type="checkbox"/>	SEA	HA	6.8554526749	
<input type="checkbox"/>	SEA	AA	6.93982738752	
<input type="checkbox"/>	SEA	CO	7.09645886862	

Query: [Index] CarriersByDepDelayTask2: Origin, DepDelay ▾				
<input type="checkbox"/>	Origin ▾	Carrier ▾	DepDelay ▾	
<input type="checkbox"/>	BOS	TZ	3.06379208506	
<input type="checkbox"/>	BOS	PA (1)	4.44716479505	
<input type="checkbox"/>	BOS	ML (1)	5.73477564103	
<input type="checkbox"/>	BOS	EV	7.20813771518	
<input type="checkbox"/>	BOS	NW	7.24518878651	
<input type="checkbox"/>	BOS	DL	7.44120228111	
<input type="checkbox"/>	BOS	XE	8.10292249047	
<input type="checkbox"/>	BOS	US	8.68794683166	
<input type="checkbox"/>	BOS	AA	8.72883304265	
<input type="checkbox"/>	BOS	EA	8.90206833904	

Problem 2.2

Strategy:

Same as problem 2.1 except replace carrier with DestAirport

Source:

[g2_2.py](#)

Result:

Output format:

Origin, ((dest1, delay), (dest2, delay)....)

```
19/06/09 15:22:03 INFO CheckpointWriter: Checkpoint for time 1560093723000 ms saved to file 's3://mudabircapstonecheckpoint/top1airportsyairport/checkpoint-1560093723000', took 18168 bytes
and 326 ms
(u'JFK', [(u'SWF', -10.5), (u'MYR', 0.0), (u'ABQ', 0.0), (u'ISP', 0.0), (u'ANC', 0.0), (u'UCA', 1.9170124481327802), (u'BGR', 3.210280373831776), (u'BQN', 3.606227610912097), (u'CHS', 4.4027105517909), (u'STT', 4.537363657461128)])
(u'MIA', [(u'SHV', 0.0), (u'BUF', 1.0), (u'SAN', 1.710382513661202), (u'SLC', 2.5371900826446283), (u'HOU', 2.912199124726477), (u'ISP', 3.647398843930636), (u'MEM', 3.7451066224751424), (u'PSE', 3.975845410628019), (u'TLH', 4.2614844746916205), (u'MCI', 4.612244897959184)])
(u'LAX', [(u'SDF', -16.0), (u'IDA', -7.0), (u'DRO', -6.0), (u'RSW', -3.0), (u'LAX', -2.0), (u'BZN', -0.7272727272727273), (u'PIH', 0.0), (u'IYK', 1.2698247440569148), (u'MFE', 1.3764705882352941), (u'MEM', 1.869798722663054)])
(u'CMI', [(u'ABI', -7.0), (u'PIT', 1.1024305555555556), (u'CVG', 1.8947616800377536), (u'DAY', 3.116235294117647), (u'STL', 3.981673306772908), (u'PIA', 4.591891891891892), (u'DFW', 5.944142746314973), (u'ATL', 6.665137614678899), (u'ORD', 8.194098143236074)])
(u'SEA', [(u'EUG', 0.0), (u'PIH', 1.0), (u'PSC', 2.6505190311418687), (u'CVG', 3.878744557801027), (u'MEM', 4.26022369800769), (u'CLE', 5.1701694915254235), (u'BLI', 5.198249133685938), (u'YKM', 5.379647749510763), (u'SNA', 5.406250794054123), (u'LIH', 5.481081081081081)])
(u'BOS', [(u'SWF', -5.0), (u'ONT', -3.0), (u'GGG', 1.0), (u'AUS', 1.2087076710435383), (u'LGA', 3.0541274274992913), (u'MSY', 3.2464678178963893), (u'LGB', 5.136176772867421), (u'OAK', 5.783210035381152), (u'MEW', 5.895637536821433), (u'BDL', 5.982704848313014)])
(u'BWI', [(u'SAV', -7.0), (u'MLB', 1.155367231638418), (u'DAB', 1.4695945945945945), (u'SRQ', 1.5884838880084522), (u'IAD', 1.7909407665505226), (u'UCA', 3.6541698546289214), (u'CHO', 3.744927536231884), (u'GSP', 4.197686645636172), (u'OAJ', 4.471111111111111), (u'SJU', 4.473430447271235)])
(u'CMH', [(u'SYR', -5.0), (u'AUS', -5.0), (u'OMA', -5.0), (u'MSN', 1.0), (u'CLE', 1.10498687664042), (u'SDF', 1.3529411764705883), (u'CAK', 3.700394218134034), (u'SLC', 3.9392857142857145), (u'MEM', 4.152021563342318), (u'IAD', 4.158103448275862)])
(u'SFO', [(u'SDF', -10.0), (u'NSO', -4.0), (u'PIH', -3.0), (u'LGA', -1.7575757575757576), (u'PIE', -1.341040462427457), (u'OAK', -0.813200498132005), (u'FAR', 0.0), (u'BNA', 2.425966447848286), (u'MEM', 3.302482299752623), (u'SCK', 4.0)])
(u'SRQ', [(u'YXW', 0.0), (u'TPA', 1.3288513253937764), (u'IAH', 1.4445574771108851), (u'MEM', 1.7029598308668077), (u'FLL', 2.0), (u'BNA', 2.0623145400593472), (u'MCO', 2.364537698870187), (u'RDW', 2.535400709882309), (u'MDW', 2.838123554674595), (u'CLT', 3.358363542206111)])
19/06/09 15:22:04 INFO JobScheduler: Finished job streaming job 1560093723000 ms.0 from job set of time 1560093723000 ms
19/06/09 15:22:04 INFO JobScheduler: Starting job streaming job 1560093723000 ms.1 from job set of time 1560093723000 ms
```

From logs:

and 322 ms

(u'JFK', [(u'SWF', -10.5), (u'MYR', 0.0), (u'ABQ', 0.0), (u'ISP', 0.0), (u'ANC', 0.0), (u'UCA', 1.9170124481327802), (u'BGR', 3.210280373831776), (u'BQN', 3.606227610912097), (u'CHS', 4.4027105517909), (u'STT', 4.537363657461128)])
(u'MIA', [(u'SHV', 0.0), (u'BUF', 1.0), (u'SAN', 1.710382513661202), (u'SLC', 2.5371900826446283), (u'HOU', 2.912199124726477), (u'ISP', 3.647398843930636), (u'MEM', 3.7451066224751424), (u'PSE', 3.975845410628019), (u'TLH', 4.2614844746916205), (u'MCI', 4.612244897959184)])
(u'LAX', [(u'SDF', -16.0), (u'IDA', -7.0), (u'DRO', -6.0), (u'RSW', -3.0), (u'LAX', -2.0), (u'BZN', -0.7272727272727273), (u'PIH', 0.0), (u'IYK', 1.2698247440569148), (u'MFE', 1.3764705882352941), (u'MEM', 1.869798722663054)])
(u'CMI', [(u'ABI', -7.0), (u'PIT', 1.1024305555555556), (u'CVG', 1.8947616800377536), (u'DAY', 3.116235294117647), (u'STL', 3.981673306772908), (u'PIA', 4.591891891891892), (u'DFW', 5.944142746314973), (u'ATL', 6.665137614678899), (u'ORD', 8.194098143236074)])
(u'SEA', [(u'EUG', 0.0), (u'PIH', 1.0), (u'PSC', 2.6505190311418687), (u'CVG', 3.878744557801027), (u'MEM', 4.26022369800769), (u'CLE', 5.1701694915254235), (u'BLI', 5.198249133685938), (u'YKM', 5.379647749510763), (u'SNA', 5.406250794054123), (u'LIH', 5.481081081081081)])

(u'BOS', [(u'SWF', -5.0), (u'ONT', -3.0), (u'GGG', 1.0), (u'AUS', 1.2087076710435383), (u'LGA', 3.0541274274992913), (u'MSY', 3.2464678178963893), (u'LGB', 5.136176772867421), (u'OAK', 5.783210035381152), (u'MDW', 5.895637536821433), (u'BDL', 5.982704848313014)]) (u'BWI', [(u'SAV', -7.0), (u'MLB', 1.155367231638418), (u'DAB', 1.4695945945945945), (u'SRQ', 1.5884838880084522), (u'IAD', 1.7909407665505226), (u'UCA', 3.6541698546289214), (u'CHO', 3.744927536231884), (u'GSP', 4.197686645636172), (u'OAJ', 4.471111111111111), (u'SJU', 4.473430447271235)]) (u'CMH', [(u'SYR', -5.0), (u'AUS', -5.0), (u'OMA', -5.0), (u'MSN', 1.0), (u'CLE', 1.10498687664042), (u'SDF', 1.3529411764705883), (u'CAK', 3.700394218134034), (u'SLC', 3.9392857142857145), (u'MEM', 4.152021563342318), (u'IAD', 4.158103448275862)]) (u'SFO', [(u'SDF', -10.0), (u'MSO', -4.0), (u'PIH', -3.0), (u'LGA', -1.7575757575757576), (u'PIE', -1.3410404624277457), (u'OAK', -0.813200498132005), (u'FAR', 0.0), (u'BNA', 2.425966447848286), (u'MEM', 3.302482299752623), (u'SCK', 4.0)]) (u'SRQ', [(u'EYW', 0.0), (u'TPA', 1.3288513253937764), (u'IAH', 1.4445574771108851), (u'MEM', 1.7029598308668077), (u'FLL', 2.0), (u'BNA', 2.0623145400593472), (u'MCO', 2.364537698870187), (u'RDU', 2.535400709882309), (u'MDW', 2.838123554674595), (u'CLT', 3.358363542206111)])

19/06/09 15:21:24 INFO JobScheduler: Finished job streaming job 1560093684000 ms.0 from job set of time 1560093684000 ms

19/06/09 15:21:24 INFO JobScheduler: Starting job streaming job 1560093684000 ms.1 from job set of time 1560093684000 ms

19/06/09 15:21:25 INFO JobScheduler: Finished job streaming job 1560093684000 ms.1 from job set of time 1560093684000 ms

19/06/09 15:21:25 INFO JobScheduler: Total delay: 1.318 s for time 1560093684000 ms (execution: 1.206 s)

19/06/09 15:21:25 INFO KafkaRDD: Removing RDD 21220 from persistence list

19/06/09 15:21:25 INFO JobGenerator: Checkpointing graph for time 1560093684000 ms

19/06/09 15:21:25 INFO DStreamGraph: Updating checkpoint data for time 1560093684000 ms

19/06/09 15:21:25 INFO DStreamGraph: Updated checkpoint data for time 1560093684000 ms

<input type="checkbox"/>	Origin	Dest	DepDelay
<input type="checkbox"/>	JFK	SWF	-10.5
<input type="checkbox"/>	JFK	ISP	0
<input type="checkbox"/>	JFK	ABQ	0
<input type="checkbox"/>	JFK	MYR	0
<input type="checkbox"/>	JFK	ANC	0
<input type="checkbox"/>	JFK	UCA	1.91701244813
<input type="checkbox"/>	JFK	BGR	3.21028037383
<input type="checkbox"/>	JFK	BQN	3.60622761091
<input type="checkbox"/>	JFK	CHS	4.40271055179
<input type="checkbox"/>	JFK	STT	4.53736365746

<input type="checkbox"/>	Origin ▾	Dest ▾	DepDelay ▾
<input type="checkbox"/>	SEA	EUG	0
<input type="checkbox"/>	SEA	PIH	1
<input type="checkbox"/>	SEA	PSC	2.65051903114
<input type="checkbox"/>	SEA	CVG	3.8787445578
<input type="checkbox"/>	SEA	MEM	4.26022369801
<input type="checkbox"/>	SEA	CLE	5.17016949153
<input type="checkbox"/>	SEA	BLI	5.19824913369
<input type="checkbox"/>	SEA	YKM	5.37964774951
<input type="checkbox"/>	SEA	SNA	5.40625079405
<input type="checkbox"/>	SEA	LIH	5.48108108108

<input type="checkbox"/>	Origin ▾	Dest ▾	DepDelay ▾
<input type="checkbox"/>	CMH	SYR	-5
<input type="checkbox"/>	CMH	OMA	-5
<input type="checkbox"/>	CMH	AUS	-5
<input type="checkbox"/>	CMH	MSN	1
<input type="checkbox"/>	CMH	CLE	1.10498687664
<input type="checkbox"/>	CMH	SDF	1.35294117647
<input type="checkbox"/>	CMH	CAK	3.70039421813
<input type="checkbox"/>	CMH	SLC	3.93928571429
<input type="checkbox"/>	CMH	MEM	4.15202156334
<input type="checkbox"/>	CMH	IAD	4.15810344828

<input type="checkbox"/>	Origin ▾	Dest ▾	DepDelay ▾
<input type="checkbox"/>	SRQ	EYW	0
<input type="checkbox"/>	SRQ	TPA	1.32885132539
<input type="checkbox"/>	SRQ	IAH	1.44455747711
<input type="checkbox"/>	SRQ	MEM	1.70295983087
<input type="checkbox"/>	SRQ	FLL	2
<input type="checkbox"/>	SRQ	BNA	2.06231454006
<input type="checkbox"/>	SRQ	MCO	2.36453769887
<input type="checkbox"/>	SRQ	RDU	2.53540070988
<input type="checkbox"/>	SRQ	MDW	2.83812355467
<input type="checkbox"/>	SRQ	CLT	3.35836354221

<input type="checkbox"/>	Origin	Dest	DepDelay
<input type="checkbox"/>	BOS	SWF	-5
<input type="checkbox"/>	BOS	ONT	-3
<input type="checkbox"/>	BOS	GGG	1
<input type="checkbox"/>	BOS	AUS	1.20870767104
<input type="checkbox"/>	BOS	LGA	3.0541274275
<input type="checkbox"/>	BOS	MSY	3.2464678179
<input type="checkbox"/>	BOS	LGB	5.13617677287
<input type="checkbox"/>	BOS	OAK	5.78321003538
<input type="checkbox"/>	BOS	MDW	5.89563753682
<input type="checkbox"/>	BOS	BDL	5.98270484831

Problem 2.4

(determine the mean arrival delay (in minutes) for a flight from X to Y)

Strategy:

1. Filter for all non cancelled flights and map (origin,dest) -> (ArrDelay)
2. Get (origin, dest) -> (ArrDelaySum,count, avgArrDelay) // See flightsDelay.updateStateByKey(updateFunction)
3. Save to DB

Source:

[g2_4.py](#)

Result:

Format: { (<origin>, <dest>), <delay value> }

```
((u'LGA', u'BOS'), 1.4838648387077622)
((u'CMI', u'ORD'), 10.14366290643663)
((u'JFK', u'LAX'), 6.635119155270517)
((u'BOS', u'LGA'), 3.7841181478417854)
((u'OKC', u'DFW'), 5.027862768428806)
((u'LAX', u'SFO'), 9.589282731105238)
((u'ATL', u'PHX'), 9.021341881513989)
((u'DFW', u'IAH'), 7.617332798592114)
((u'MSP', u'ATL'), 6.737007973674219)
((u'IND', u'CMH'), 2.8911367050575865)
```

```

19/06/09 16:03:27 INFO CheckpointWriter: Submitted checkpoint of time 1560096207000 ms to writer queue
19/06/09 16:03:27 INFO CheckpointWriter: Saving checkpoint for time 1560096207000 ms to file 's3://mudabircapstonecheckpoint/meanDelayBetweenAandB/checkpoint-1560096207000'
19/06/09 16:03:27 INFO CheckpointWriter: Deleting s3://mudabircapstonecheckpoint/meanDelayBetweenAandB/checkpoint-1560096192000.bk
19/06/09 16:03:27 INFO CheckpointWriter: Checkpoint for time 1560096207000 ms saved to file 's3://mudabircapstonecheckpoint/meanDelayBetweenAandB/checkpoint-1560096207000', took 17310 bytes and 319 ms
((u'LGA', u'BOS'), 1.4838648387077622)
((u'CMT', u'ORD'), 10.14366290643663)
((u'JFK', u'LAX'), 6.635119155270517)
((u'BOS', u'LGA'), 3.7841181478417854)
((u'OKC', u'DFW'), 5.027862768428806)
((u'LAX', u'SFO'), 9.589282731105238)
((u'DFW', u'IAH'), 7.617332798592114)
((u'ATL', u'PHX'), 9.021341881513989)
((u'MSP', u'ATL'), 6.737007973674219)
((u'IND', u'CMH'), 2.8911367050575865)
19/06/09 16:03:27 INFO JobScheduler: Finished job streaming job 1560096207000 ms.0 from job set of time 1560096207000 ms
19/06/09 16:03:27 INFO JobScheduler: Starting job streaming job 1560096207000 ms.1 from job set of time 1560096207000 ms
19/06/09 16:03:27 INFO JobScheduler: Finished job streaming job 1560096207000 ms.1 from job set of time 1560096207000 ms
19/06/09 16:03:27 INFO JobScheduler: Total delay: 0.655 s for time 1560096207000 ms (execution: 0.539 s)
19/06/09 16:03:27 INFO KafkaRDD: Removing RDD 14320 from persistence list
19/06/09 16:03:27 INFO JobGenerator: Checkpointing graph for time 1560096207000 ms
19/06/09 16:03:27 INFO DStreamGraph: Updating checkpoint data for time 1560096207000 ms
19/06/09 16:03:27 INFO DStreamGraph: Updated checkpoint data for time 1560096207000 ms

```

Arrival delay of some origin-dest pairs from dynamoDB

Scan: [\[Table\] MeanDelayBetweenAandBTask2: AtoB](#) ▾

<input type="checkbox"/>	AtoB ⓘ ▲	ArrDelay ▼
<input type="checkbox"/>	(u'ATL', u'PHX')	9.02134188151
<input type="checkbox"/>	(u'BOS', u'LGA')	3.78411814784
<input type="checkbox"/>	(u'CMT', u'ORD')	10.1436629064
<input type="checkbox"/>	(u'DFW', u'IAH')	7.61733279859
<input type="checkbox"/>	(u'IND', u'CMH')	2.89113670506
<input type="checkbox"/>	(u'JFK', u'LAX')	6.63511915527
<input type="checkbox"/>	(u'LAX', u'SFO')	9.58928273111
<input type="checkbox"/>	(u'LGA', u'BOS')	1.48386483871
<input type="checkbox"/>	(u'MSP', u'ATL')	6.73700797367
<input type="checkbox"/>	(u'OKC', u'DFW')	5.02786276843

Problem 3.1

(copy pasting my results from Task1 for this problem)

Strategy:

- 1) Collected ranking of all airports into a file using Spark
- 2) Used powerlaw library to plot the CCDF of the power law distribution and the lognormal distribution of the data collected in 1)

Source:


[g3_1.py](#)

[pythonplot.py](#)

Result:

[g3_1.log](#) (Ranking of airports with their count)

Our data resembles the log Normal fit better than the powerlaw fit. We can conclude that the popularity of airports doesn't follow Zipf distribution



Problem 3.2

Strategy:

1. Filter non cancelled flights

2. Create RDD for all flights which fly before noon and dest as Y. key is (date, Y) -> (flight info)
3. Create RDD for all flights which fly after noon with origin as Y and departure date subtracted by 2 days
4. Join is done on (date,Y) as key. This gives all flights landing in Y before noon and all flights departing from Y two days later
5. For X->Y->Z route, filter the flight combo with minimum arrival delay
6. Write data to DynamoDB. Since a huge set of data was to be written, increased the write capacity of DynamoDB to 1500 to transfer the data faster.

Source:

[g3_2.py](#)

Result:

```
19/06/09 16:15:51 INFO JobScheduler: Starting job streaming job 1560096951000 ms.0 from job set of time 1560096951000 ms
19/06/09 16:15:51 INFO DStreamGraph: Updated checkpoint data for time 1560096951000 ms
19/06/09 16:15:51 INFO CheckpointWriter: Submitted checkpoint of time 1560096951000 ms to writer queue
19/06/09 16:15:51 INFO CheckpointWriter: Saving checkpoint for time 1560096951000 ms to file 's3://mudabircapstonecheckpoint/bestFlights/checkpoint-1560096951000'
*****
(('2008-04-03', 'BOS', 'ATL', 'LAX'), (((u'BOS', u'ATL', u'FL', u'270', u'0600', 7.0), (u'ATL', u'LAX', u'FL', u'40', u'1852', -2.0)), 5.0))
*****
(('2008-09-09', 'JAX', 'DFW', 'CRP'), (((u'JAX', u'DFW', u'AA', u'845', u'0725', 1.0), (u'DFW', u'CRP', u'MQ', u'3627', u'1645', -7.0)), -6.0))
*****
(('2008-05-16', 'LAX', 'MIA', 'LAX'), (((u'LAX', u'MIA', u'AA', u'280', u'0820', 10.0), (u'MIA', u'LAX', u'AA', u'456', u'1930', -19.0)), -9.0))
*****
(('2008-09-07', 'PHX', 'JFK', 'MSP'), (((u'PHX', u'JFK', u'B6', u'178', u'1130', -25.0), (u'JFK', u'MSP', u'NW', u'609', u'1750', -17.0)), -42.0))
*****
(('2008-01-01', 'LAX', 'ORD', 'JFK'), (((u'LAX', u'ORD', u'UA', u'944', u'0705', 1.0), (u'ORD', u'JFK', u'B6', u'918', u'1900', -7.0)), -6.0))
*****
(('2008-07-12', 'LAX', 'SFO', 'PHX'), (((u'LAX', u'SFO', u'WN', u'3534', u'0650', -13.0), (u'SFO', u'PHX', u'US', u'412', u'1925', -19.0)), -32.0))
*****
(('2008-04-01', 'SLC', 'BFL', 'LAX'), (((u'SLC', u'BFL', u'OO', u'3755', u'1100', 12.0), (u'BFL', u'LAX', u'OO', u'5429', u'1455', 6.0)), 18.0))
*****
(('2008-01-24', 'DFW', 'STL', 'ORD'), (((u'DFW', u'STL', u'AA', u'1336', u'0705', -14.0), (u'STL', u'ORD', u'AA', u'2245', u'1655', -5.0)), -19.0))
*****
(('2008-06-10', 'DFW', 'ORD', 'DFW'), (((u'DFW', u'ORD', u'UA', u'1104', u'0700', -21.0), (u'ORD', u'DFW', u'AA', u'2341', u'1645', -10.0)), -31.0))
*****
(('2008-03-04', 'CMI', 'ORD', 'LAX'), (((u'CMI', u'ORD', u'MQ', u'4278', u'0710', -14.0), (u'ORD', u'LAX', u'AA', u'607', u'1950', -24.0)), -38.0))
19/06/09 16:15:51 INFO JobScheduler: Finished job streaming job 1560096951000 ms.0 from job set of time 1560096951000 ms
19/06/09 16:15:51 INFO JobScheduler: Total delay: 0.512 s for time 1560096951000 ms (execution: 0.350 s)
19/06/09 16:15:51 INFO KafkaRDD: Removing RDD 2502 from persistence list
19/06/09 16:15:51 INFO JobGenerator: Checkpointing graph for time 1560096951000 ms
19/06/09 16:15:51 INFO DStreamGraph: Updating checkpoint data for time 1560096951000 ms
19/06/09 16:15:51 INFO DStreamGraph: Updated checkpoint data for time 1560096951000 ms
19/06/09 16:15:51 INFO CheckpointWriter: Deleting s3://mudabircapstonecheckpoint/bestFlights/checkpoint-1560096951000.bk
```

Format: ((startdate,X,Y,Z), ((XYdetails,YZdetails),totalArrivalDelay))

((('2008-04-03', 'BOS', 'ATL', 'LAX'), (((u'BOS', u'ATL', u'FL', u'270', u'0600', 7.0), (u'ATL', u'LAX', u'FL', u'40', u'1852', -2.0)), 5.0))

((('2008-09-09', 'JAX', 'DFW', 'CRP'), (((u'JAX', u'DFW', u'AA', u'845', u'0725', 1.0), (u'DFW', u'CRP', u'MQ', u'3627', u'1645', -7.0)), -6.0))

((('2008-05-16', 'LAX', 'MIA', 'LAX'), (((u'LAX', u'MIA', u'AA', u'280', u'0820', 10.0), (u'MIA', u'LAX', u'AA', u'456', u'1930', -19.0)), -9.0))

```

*****__****
(('2008-09-07', 'PHX', 'JFK', 'MSP'), (((u'PHX', u'JFK', u'B6', u'178', u'1130', -25.0), (u'JFK',
u'MSP', u'NW', u'609', u'1750', -17.0)), -42.0))
*****__****
(('2008-01-01', 'LAX', 'ORD', 'JFK'), (((u'LAX', u'ORD', u'UA', u'944', u'0705', 1.0), (u'ORD',
u'JFK', u'B6', u'918', u'1900', -7.0)), -6.0))
*****__****
(('2008-07-12', 'LAX', 'SFO', 'PHX'), (((u'LAX', u'SFO', u'WN', u'3534', u'0650', -13.0), (u'SFO',
u'PHX', u'US', u'412', u'1925', -19.0)), -32.0))
*****__****
(('2008-04-01', 'SLC', 'BFL', 'LAX'), (((u'SLC', u'BFL', u'OO', u'3755', u'1100', 12.0), (u'BFL',
u'LAX', u'OO', u'5429', u'1455', 6.0)), 18.0))
*****__****
(('2008-01-24', 'DFW', 'STL', 'ORD'), (((u'DFW', u'STL', u'AA', u'1336', u'0705', -14.0), (u'STL',
u'ORD', u'AA', u'2245', u'1655', -5.0)), -19.0))
*****__****
(('2008-06-10', 'DFW', 'ORD', 'DFW'), (((u'DFW', u'ORD', u'UA', u'1104', u'0700', -21.0), (u'ORD',
u'DFW', u'AA', u'2341', u'1645', -10.0)), -31.0))
*****__****
(('2008-03-04', 'CMI', 'ORD', 'LAX'), (((u'CMI', u'ORD', u'MQ', u'4278', u'0710', -14.0), (u'ORD',
u'LAX', u'AA', u'607', u'1950', -24.0)), -38.0))

```

Snapshot from DynamoDB

Scan: [Table] BestArrivalTimeFinalTask2: XYZ, StartDate ▾					Viewing 1 to 10 items
<input type="checkbox"/>	XYZ ▾	StartDate ▾	ArrDelay ▾	info ⓘ	
<input type="checkbox"/>	BOS-ATL-LAX	2008-04-03	5	((u'BOS', u'ATL', u'FL', u'270', u'0600', 7.0), (u'ATL', u'LAX', u'FL', u'40', u'1852', -2.0))	
<input type="checkbox"/>	CMI-ORD-LAX	2008-03-04	-38	((u'CMI', u'ORD', u'MQ', u'4278', u'0710', -14.0), (u'ORD', u'LAX', u'AA', u'607', u'1950', -24.0))	
<input type="checkbox"/>	DFW-ORD-DFW	2008-06-10	-31	((u'DFW', u'ORD', u'UA', u'1104', u'0700', -21.0), (u'ORD', u'DFW', u'AA', u'2341', u'1645', -10.0))	
<input type="checkbox"/>	DFW-STL-ORD	2008-01-24	-19	((u'DFW', u'STL', u'AA', u'1336', u'0705', -14.0), (u'STL', u'ORD', u'AA', u'2245', u'1655', -5.0))	
<input type="checkbox"/>	JAX-DFW-CRP	2008-09-09	-6	((u'JAX', u'DFW', u'AA', u'845', u'0725', 1.0), (u'DFW', u'CRP', u'MQ', u'3627', u'1645', -7.0))	
<input type="checkbox"/>	LAX-MIA-LAX	2008-05-16	-9	((u'LAX', u'MIA', u'AA', u'280', u'0820', 10.0), (u'MIA', u'LAX', u'AA', u'456', u'1930', -19.0))	
<input type="checkbox"/>	LAX-ORD-JFK	2008-01-01	-6	((u'LAX', u'ORD', u'UA', u'944', u'0705', 1.0), (u'ORD', u'JFK', u'B6', u'918', u'1900', -7.0))	
<input type="checkbox"/>	LAX-SFO-PHX	2008-07-12	-32	((u'LAX', u'SFO', u'WN', u'3534', u'0650', -13.0), (u'SFO', u'PHX', u'US', u'412', u'1925', -19.0))	
<input type="checkbox"/>	PHX-JFK-MSP	2008-09-07	-42	((u'PHX', u'JFK', u'B6', u'178', u'1130', -25.0), (u'JFK', u'MSP', u'NW', u'609', u'1750', -17.0))	
<input type="checkbox"/>	SLC-BFL-LAX	2008-04-01	18	((u'SLC', u'BFL', u'OO', u'3755', u'1100', 12.0), (u'BFL', u'LAX', u'OO', u'5429', u'1455', 6.0))	

Optimizations

- Filtered only necessary fields from all the csv reducing the total data size to 4.5GB from 15 GB
- In the kafka producer side, increased batch size and linger_ms to dispatch larger amounts of data at a time to the kafka cluster without waiting or ack. Also make ack=True to make sure no data is dropped/lost while writing to kafka topic

- Made replication factor =3 to make the kafka cluster reliable. Kept the number of partitions to 1. Increasing the number of partitions seemed to make the read from consumer side slower

Streaming versus Batch processing

- For the given data, batch processing seemed a better option. Since we had to compute top 10 airports/carriers etc based on entire data, doing stream processing as the data arrives seems like an overkill.
- However, stream processing is useful in analysing live how the parameters are changing over time.

Comparing Stacks

Since I used Spark for task1 and Spark Streaming for task2, the stacks were more or less similar. Using hadoop Map-reduce for task 1 would have made the process very slow.
